

AIStudio Benchmark Audit — ESEF Banks Corpus

Run C · Full Corpus, Blended Language (`esef_blended`) · 2026-06-26

Run ID: benchmark_esef_banks_gemma3-27b_top_k10_2026-06-26_2222 **Auditor:** Claude (Opus) · audited against ESEF primary filings (filings.xbrl.org) **Corpus:** esef_banks · 12 firms · scope: none (full corpus) · 5 EN-filing + 7 non-EN-filing **Config:** Top-K=10 · T=0.3 · α =0.5 · gemma3:27b · entity-filter:auto · keywords:on **Questions:** default set · 12 questions **Companion:** Run D (`esef_english`) — the same corpus restricted to the 5 English filers with matched questions

Companion reading — this report is part of the canonical benchmark suite. Read it alongside [TUTORIAL.md](#) §5 (how to read a benchmark) and Annex 3 (the language ceiling), and see the suite synthesis: [BENCH - Canonical Suite - README and Synthesis](#). The reports carry the data and the audit; the *method* lives in the Tutorial, the *mechanism* in its Annexes.

Executive Summary

Mechanical: 8/12 (6 · 6 · 0) · **Calibrated:** ~8–9/12 · 0 fabrications · **Avg latency** 62.4s

Run C is the **blended-language baseline**: all 12 ESEF banks, default questions, no scope filter — the number a user sees querying the full European corpus without knowing which filings are in English. Five firms file in English (ING, Barclays, HSBC, Standard Chartered, BBVA); seven file in their domestic language (Nordea/SEB Swedish, Erste German, KBC Dutch, BNP/Société Générale French, UniCredit Italian). The run's defining feature: **every soft spot falls on a non-English firm, and not one is a fabrication**. Where the system can't ground an answer in a foreign-language filing, it returns a sparse-but-correct answer, drifts to an adjacent section, or — in the cleanest case — explicitly abstains. It never invents.

The mechanical 8/12 both *under-* and *over-*states quality in different places, which is why it is audited rather than trusted: three ambers are correct answers with sparse attribution (read up), one GREEN hides an abstention (read down).

Results at a Glance

#	ID	Firm (filing lang)	Mech	Calibrated	One-line
1	cet1_ing_group	ING (en)			Correct CET1 13.1%/buffer; amber is low density only

#	ID	Firm (filing lang)	Mech	Calibrated	One-line
2	tier1_capital_barclays	Barclays (en)	•	•	Full capital stack, correct figures
3	leverage_ratio_hsbc	HSBC (en)	•		5.3% + UK 3.25% + buffers — primary-verified correct; "regulatory" keyword artifact
4	cet1_cross_firm_en	ING/HSBC/ Barclays (en)	•	×	Three-firm CET1, correct, no contamination
5	nordea_leverage_ratio	Nordea (sv)	•	▀ abstain	GREEN hides an abstention: "sources do not directly address this"
6	kbc_nii_sensitivity	KBC (nl)	•	▀ wrong-topic	Answered ESG/climate credit risk, not the asked NII sensitivity
7	climate_risk_bbva	BBVA (en)	•	•	TCFD/GFANZ transition plan, correct; low density
8	climate_standard_chartered	Std Chartered (en)	•	•	Net-zero ops by 2025, scenario analysis, correct
9	digital_ing	ING (en)	•	•	Four-enabler tech strategy, correct
10	digital_nordea	Nordea (sv)	•	—	Correct (top-3 digital, app awards); "technology" keyword artifact
11	barclays_holdco_sub	Barclays (en)			Holdco/sub structure, FHC status, correct
12	erste_whitespace_collapse	Erste (de)		—	Sparse but real risk-profile answer; "framework" keyword artifact

Notable Per-Question Findings

The language pattern (the run's thesis)

Sort the soft spots by filing language and the result is unambiguous: **all four English firms with a flagged question (Q1 ING, Q3 HSBC, Q7 BBVA) amber only on keyword/density technicalities while answering correctly; the genuine quality problems are on non-English firms** — Q5 (Nordea, sv), Q6 (KBC, nl), Q10 (Nordea, sv), Q12 (Erste, de). This is the language ceiling of **Annex 3** at corpus scale: retrieval quality is not language-neutral, and a question in English against a filing written in Swedish/Dutch/German retrieves worse.

Q5 › `nordea_leverage_ratio` — the GREEN that hides an abstention (read down)

Scored GREEN 1.0, but the answer is: *"The available sources do not directly address this question..."* — an honest abstention on a Swedish filing, scored as a pass because the citation and topical keywords matched. This is the mirror of B's false-GREENs: here the mechanical score over-credits a non-answer. The right read is "could not ground the specific leverage figure across the language gap" — a known non-English limitation, correctly handled by abstaining rather than inventing.

Q6 `kbc_nii_sensitivity` — wrong-section retrieval (the one real miss)

The ask is net-interest-income sensitivity and interest-rate-risk management. The answer describes KBC's **ESG/climate credit-risk** framework instead — a wrong-section retrieval on a Dutch filing. This is the genuine miss of the run: not a fabrication, but the wrong part of the document surfaced because the English query matched the wrong Dutch passage.

Q3 `leverage_ratio_hsb` — primary-verified, amber is a keyword artifact (read up)

HSBC 5.3% (from 5.6%), UK minimum 3.25%, 0.9% buffer (0.7% ALRB + 0.2% CCLB), exceeded throughout 2025 — every figure verified against HSBC's Pillar 3 disclosure. The amber is solely the missing literal token "regulatory"; the answer is fully correct. Read up.

Synthesis

Run C is the **honest blended baseline** for the European corpus: 8/12 mechanical, ~8–9/12 calibrated, on a 12-firm portfolio where seven firms file in a non-English language. Its value is in *where* the soft spots fall — exclusively on the non-English firms, and exclusively as sparse-attribution, wrong-section, or honest-abstention, **never fabrication**. The system degrades gracefully across the language gap: it says less, points at less, or declines — it does not invent a plausible foreign-sourced number.

The cross-firm English questions (Q4, and the single-firm English Q1/Q2/Q3/Q7/Q8/Q9/Q11) are correct, confirming entity isolation holds at 12-firm scale with no contamination. The European-specific frontier is the language one, documented in **Annex 3**: the glossary bridges acronym ↔ full-form within English but does not translate, so an English question against a Swedish or Dutch filing retrieves worse. Run D isolates the clean half of this corpus to show what the system does when language is controlled.

Claims this run supports: - Entity isolation holds at 12-firm scale; zero cross-firm contamination, zero fabrication - The system degrades *gracefully* on non-English filings — sparse, wrong-section, or honest abstention (Q5/Q6/Q10/Q12), never invention - English-firm answers are correct and primary-verifiable (HSBC leverage, ING CET1, three-firm cross-CET1) - The frontier is language, not retrieval mechanics (Annex 3); Run D tests the language-controlled cut directly

*Report basis: benchmark_esef_banks_gemma3-27b_top_k10_2026-06-26_2222.{md,json} · audited
against ESEF primary filings + HSBC Pillar 3 · filing languages verified by chunk-content read ·
2026-06-26*

* * * * *