

Architecture Decisions

Version: 1.3.0 | Updated: 2026-07-19

Key technical choices made in this project, with rationale. This document is intended for anyone reviewing the codebase who wants to understand not just what was built, but why.

1. Local-First Design (Ollama, No External API)

Decision: All LLM inference and embedding runs locally via Ollama. No OpenAI, Anthropic, or other cloud API dependency in the core stack.

Why: - Forces genuine understanding of model behavior at the infrastructure level — latency, memory pressure, model selection tradeoffs - Eliminates API cost as a variable during experimentation - Relevant to financial services contexts where data residency and network egress are real constraints, not afterthoughts - Makes the system runnable in air-gapped or restricted environments

Tradeoff: Inference quality is lower than frontier models. This is a deliberate constraint, not a limitation. The goal is substrate knowledge, not benchmark performance.

Would change for production: Add LiteLLM as a provider abstraction layer. Local inference for development and sensitive data; cloud inference (with appropriate data handling) for production quality. The architecture is designed to make this substitution straightforward — Ollama client is isolated in `ollama_client.py`.

2. Qdrant as Vector Store (replaced ChromaDB)

Decision: Qdrant 1.17.0, running as a separate process on port 6333. ChromaDB was the original choice and was replaced in production.

Why the switch: ChromaDB crashed at 32,285 chunks during the SEC 10-K corpus ingest. Not a configuration issue — a hard failure at scale. Qdrant was stable at 100,659 chunks with zero failures across multiple full re-ingests.

Why Qdrant specifically: - Written in Rust — near-zero GC overhead vs ChromaDB's Python-based memory model - Native metadata filtering via `Filter / FieldCondition` — firm and year filters run at the vector search layer, not post-hoc on results - Production upgrade path: sharding, replication, quantization, gRPC (port 6334) - Separate process model — independently restartable, independently scalable, maps directly to a Docker container in cloud deployment

The four-process pattern is deliberate: Browser → FastAPI → Qdrant → Ollama. Each process independently restartable. The separation that looks like operational overhead in a personal lab is the same separation that enables horizontal scaling in production — add Qdrant nodes, add uvicorn workers, point at shared Qdrant cluster. No architecture change required.

Tradeoff: Qdrant requires a separate process and persistent storage directory (`~/qdrant_storage/`). ChromaDB ran in-process with zero operational overhead. The stability gain at scale makes this tradeoff non-negotiable.

Would change for production: Same Qdrant, configured for replication and S3-backed snapshots. Alternatively Qdrant Cloud for managed hosting. The `qdrant_store.py` adapter makes the substitution a configuration change.

3. Embedding Model: `nomic-embed-text`

Decision: `nomic-embed-text` as the default embedding model (768 dimensions, cosine similarity).

Why: - Best CPU/unified memory performance among locally available models at time of selection — low latency without GPU - Produces meaningful semantic representations for the document types in scope - Passes embedding arithmetic tests (King – Man + Woman ≈ Queen), a practical signal of semantic quality

Considered: `bge-large-en` — approximately 5% quality improvement on retrieval benchmarks, but roughly 2x slower. The tradeoff didn't justify it.

Would change for production: Benchmark against the specific corpus and query patterns in scope. Embedding model selection is highly workload-dependent.

4. Citation Logic Embedded in API, Not Modularized

Decision: Citation extraction and formatting lives in `api.py` rather than as a separate module.

Why: - Citation logic is tightly coupled to the response pipeline — it operates on LLM output immediately before it's returned - Separating it adds indirection without architectural benefit at this scale - Simpler deployment: one file to update, no import chain to maintain

Tradeoff: `api.py` is longer as a result.

Would change for production: With multiple citation formats or output channels, a dedicated citation module with a clean interface would be worth the overhead.

5. Semantic Chunking Over Fixed-Size Chunking

Decision: Boundary-aware semantic chunking. The previous implementation used fixed character counts.

Why: - Fixed chunking fragments natural language units — a sentence split across two chunks degrades retrieval significantly - Document structure carries semantic information: section headers, paragraph breaks, sentence boundaries are meaningful signals - Measured improvement: ~25–40% accuracy improvement on document queries

Implementation: Three-tier fallback — semantic boundary detection → sentence splitting → character fallback.

Tradeoff: Variable-length chunks create some unpredictability in context window usage. Handled via explicit context window management in the retriever.

6. Metadata Filtering at the Vector Layer

Decision: Firm and year filters are applied as Qdrant `FieldCondition` filters at query time, not as post-hoc filtering on results.

Why: - Post-hoc filtering wastes retrieval budget — you retrieve top-K then discard most of them. Vector-layer filtering retrieves top-K from the filtered subset. - Qdrant's native Filter API has zero measurable latency overhead. - Metadata (firm, year) is parsed from filenames at ingest time (`_parse_firm_year()` in `pipeline.py`) and stored as Qdrant payload fields. No manual tagging required.

Pattern: `Goldman_Sachs_10K_2026-02-25.htm` → `firm=Goldman Sachs, year=2026` parsed automatically. Filter adds zero latency. Filter is optional — omitting it runs cross-corpus retrieval.

7. Page-Aware PDF Chunking via pdfplumber (*implemented Beta*)

Decision: pdfplumber as primary PDF extractor, inserting `[PAGE_N]` markers at each page boundary. pypdf retained as fallback.

Why pdfplumber over pypdf: - pypdf extracts flat text with no page boundary information - pdfplumber provides page-by-page extraction with reliable boundary detection - `[PAGE_N]` markers flow through the chunking pipeline into Qdrant payload (`page` field) and `chunk_id` (`filename::page-N::chunk-M`)

Pipeline: `loaders._extract_pdf()` → `[PAGE_N]` markers in text → `pipeline.py` extracts page numbers → stored in Qdrant payload → surfaced in `RetrievedDoc.page` → rendered in UI citation footnotes with `Open ↗` link.

Null pages are expected for: non-PDF files (PPTX, DOCX), scanned PDFs with no text layer, and continuation chunks that start mid-page.

PDF viewer — click citation → scroll to page is tracked as Source Dive (pre-Jessica gate item). Page numbers are ready in the backend; the viewer frontend is the remaining work.

8. Conversation Memory: Sliding Window of 10 Turns

Decision: Last 10 conversation turns maintained as context for follow-up questions.

Why: - Covers most realistic conversational patterns without excessive context window consumption - Stateless between sessions by design — no persistence layer required

Tradeoff: Long research sessions lose early context. Vector-based memory store (storing conversation summaries as embeddings) would be more appropriate for persistent multi-session memory.

9. JSONL Baseline for Deterministic Testing

Decision: A JSONL retrieval path exists alongside the Qdrant path, returning deterministic results without embedding inference.

Why: - Embedding models are nondeterministic under some conditions and slow in CI - Allows testing the full API response pipeline without Ollama or Qdrant running - Unit tests use the JSONL path; integration tests use Qdrant

10. Single HTML File Frontend

Decision: The entire UI is `front_end/rag_studio.html` — one file, ~1,900 lines, no build step.

Why: - Zero build toolchain. Clone and open. Works immediately. - No npm, no webpack, no node_modules. Nothing to break on a new machine. - Consistent with the "runnable in under 30 minutes" QUICKSTART promise.

Current UI features (v2.0.1): Corpus management (create, rename, delete, stats, inspect), file upload with ingest progress, chat with inline citations and inline code rendering (backtick and fenced block), corpus/model change separators, auto-linkify of URLs and file paths, clickable citation filenames, corpus selector with alphabetical sort, Retrieval Mix slider (hybrid_alpha, Literal↔Conceptual), query settings sidebar, system prompt hot-reload via `/prompt/reload`.

Tradeoff: A single ~2,200-line file is harder to navigate than a componentized React app. The tradeoff is deliberate — operational simplicity over developer ergonomics at this scale.

Would change for production: React with component library for multi-user Beta. The `ui_architecture.md` doc describes the target left-bar layout.

15. Manifest-Driven Command Installation

Decision: A single manifest is the source of truth for all command routing, deployment, and installation. The installer reads the manifest to register commands — no hardcoded paths anywhere else.

Why: The previous approach used one-shot install scripts that couldn't be re-run when new commands were added, so adding a command meant editing shell configuration by hand — fragile, error-prone, non-reproducible.

The manifest-driven approach means: - A single new command can be installed by looking up its path in the manifest - The active command set can be verified against the manifest at any time - A fresh clone on a new machine installs every command from the same manifest, guaranteeing consistency - The manifest makes the command registry explicit and auditable

Tradeoff: Adding a new command now requires a manifest update before deployment. This is by design — the manifest is the deployment contract, not an afterthought.

Production Hardening (cross-cutting)

What the system would need before any network exposure:

- **Authentication:** The API has no auth. Fine for localhost; needs API key validation before any network exposure.
 - **Observability:** LLM latency, retrieval score distributions, and refusal rates are the three most important things to instrument first.
 - **Async ingestion:** Currently synchronous and blocking. Needs a queue-backed async pipeline at meaningful document volumes.
 - **Vector store backup:** Qdrant storage (`~/qdrant_storage/`) is a derived artifact but expensive to regenerate. Treat it like a database — back it up. In production: S3-backed Qdrant snapshots on a schedule.
-

11. CrossEncoder Reranker: Two-Stage Retrieval

Decision: CrossEncoder `ms-marco-MiniLM-L-6-v2` as a reranking pass after vector search, before prompt assembly.

Why: The bi-encoder (nomic-embed-text) compresses query and chunk independently into fixed vectors. Compression loses nuance — "AI governance committee" and "Firmwide Artificial Intelligence Risk and Controls Committee" are semantically equivalent but vector-distant. This is vocabulary mismatch.

The CrossEncoder reads query + chunk concatenated as a single input, with full attention across both. It scores relevance directly rather than approximating via cosine distance. Vocabulary mismatch largely disappears.

Two-stage architecture: - Stage 1: Qdrant HNSW vector search — fast, retrieves top-K candidates - Stage 2: CrossEncoder reranker — slower, reorders by true relevance

Model selection: `ms-marco-MiniLM-L-6-v2` - 22M parameters, 90MB — loads in <1s on Apple Silicon - Fine-tuned on MS MARCO (500K passage relevance pairs) — directly applicable to "rank these passages for this query" - Best latency/quality tradeoff in the MiniLM family - Knowledge-distilled from BERT-large — 6 transformer layers, preserves most quality at fraction of the depth

Implementation: `rag_core.py` — lazy load with graceful fallback if `sentence-transformers` unavailable. Zero impact on existing behavior if reranker fails to load.

Tradeoff: ~1-2s additional latency per query for CrossEncoder inference on top-K candidates. Acceptable at current corpus sizes.

12. Atomic `--force` Ingest

Decision: `--force` flag in the ingest pipeline forces reprocessing of all files regardless of whether they appear unchanged in the Qdrant manifest.

Why: Without `--force`, files already in Qdrant are skipped (MD5-based skip logic). When chunk format changes or a corpus needs a clean rebuild, `--force` bypasses the skip check and re-embeds everything.

Clarification on atomic wipe: `--force` does NOT wipe the Qdrant collection itself — it only bypasses the per-file skip check. A full corpus wipe (delete collection + re-ingest from scratch) is triggered separately via `DELETE /corpus/{name}?confirm=yes` in `api.py`, or by the `ais_start` help corpus auto-ingest which uses `--force` after collection existence check.

Implementation: `pipeline.py ingest_corpus()` — when `force=True`, `abs_path` in `qdrant_source_paths` check is bypassed. Files are re-chunked, re-embedded, and upserted regardless of prior state. Qdrant's upsert is idempotent — duplicate chunk IDs are overwritten cleanly.

Also used by: `scripts/start.sh` help corpus background refresh on every start.

13. Session Continuity via Structured State Packets

Decision: Operator workflow uses a PACKET/BUNDLE protocol to maintain full project context across AI-assisted development sessions. Each session ends with a structured markdown PACKET (produced by an internal session-state tool) capturing repo state, active pipeline, decisions made, and session summary. The PACKET is bundled with supporting documents into a zip archive attached to the next session.

Why: LLM context windows are finite and session-scoped. Without an explicit continuity mechanism, each session starts cold — the assistant has no memory of prior decisions, naming conventions, or work in progress. The PACKET protocol treats session state as a first-class artifact: structured, versioned, and portable.

The result is an effectively unbounded working memory across sessions. Complex multi-week development work — with branching decisions, evolving standards, and accumulated conventions — can be resumed with full fidelity. The protocol is also self-referential: the PACKET contains the rules for generating the next PACKET.

Pattern: a structured PACKET.md is generated, reviewed in session, then bundled with supporting docs and attached to the next thread — full context restored in under two minutes.

Broader relevance: This pattern is a prototype for the urCrew personal operating system concept — where structured state packets manage continuity across multiple parallel life and work threads, each with its own agent context. The PACKET/BUNDLE workflow is Phase 0 of that system, built and validated here.

Tradeoff: Requires discipline at session end. The session-end protocol codifies this as a non-negotiable checklist.

14. REST API Design and Publication

Decision: AIStudio exposes a REST API via FastAPI on port 8000. All UI interactions go through this API — the frontend is a pure client with no direct access to the filesystem, Qdrant, or Ollama.

Current endpoints (Beta): - `POST /ask` — RAG query with corpus, model, top-K, temperature, hybrid_alpha, keywords, conversation history - `POST /corpus/{name}/ingest` — trigger ingest of uploaded files - `GET /corpus/{name}/ingest-status` — streaming ingest progress (files, chunks, elapsed) - `GET /corpora` — list available corpora with metadata - `POST /corpus/create` — create new corpus directory structure - `POST /corpus/{name}/rename` — rename corpus directory + Qdrant collection + trigger re-ingest - `DELETE /corpus/{name}` — move corpus to macOS Trash (recoverable) - `DELETE /corpus/{name}/file/{filename}` — move file to corpus trash - `GET /corpus/{name}/info` — file list, chunk count, size - `POST /corpus/{name}/upload` — upload files to corpus - `POST /prewarm` — warm Ollama model into memory before first query - `GET /source` — serve source documents for citation Open ↗ links - `POST /prompt/reload` — hot-swap system prompt without restart (clears memoized `_SYSTEM_PROMPT_BASE`) - `GET /health` — liveness check (used by start.sh health-check poll) - `GET /models` — list available Ollama models

Why a clean API separation matters: The single-page frontend communicates exclusively via `fetch()` calls to this API. This means AIStudio's RAG capabilities are immediately consumable by any HTTP client — scripts, notebooks, other applications, or future agent orchestration layers. No SDK required.

Planned: Full `API_DOC.md` with request/response schemas and error codes published alongside the API. Swagger/OpenAPI auto-docs via FastAPI's built-in support (`/docs` endpoint). This is a v2.0 item — the API surface is stable at Beta but not yet formally documented.

16. Hybrid Retrieval: BM25 + Vector Score Combination (M2.A)

Decision: Hybrid retrieval combining Qdrant vector search with BM25 sparse retrieval, score-combined via a weighted `hybrid_alpha` parameter. Exposed as a per-query slider ("Retrieval Mix") in the UI and as a `--alpha` flag in the benchmark harness.

Why: Pure vector retrieval (semantic similarity) retrieves conceptually related chunks but can miss exact-match keyword queries — e.g. a query for "Goldman Sachs CET1 ratio 2024" may return semantically related risk content instead of the specific number. BM25 is token-based and excels at exact-match recall. Combining both captures the strengths of each.

Empirical finding: $\alpha=0.5$ (equal blend) + $K=10$ achieves 14/14 on the demo corpus benchmark. Pure vector ($\alpha=0.0$, $K=5$) achieves 13/14 — the two failures were both list-format answers where the model dropped citations. At $\alpha=0.5$, retrieval precision is higher and citation compliance improves.

Implementation: `scoring.py` — pure functions `normalize_minmax()` and `combine_hybrid()`. `qdrant_store.py` v1.1.0 adds `query_bm25()` and `_ensure_text_index()`. `rag_core.py` v1.5.0 accepts `hybrid_alpha` kwarg. `api.py` v1.6.0 exposes `hybrid_alpha` on `AskRequest`. UI slider: Literal=left ($\alpha=1.0$ BM25), Conceptual=right ($\alpha=0.0$ vector) — display shows actual α sent.

UI convention note: The slider direction inverts the raw parameter — `hybrid_alpha: 1 - parseFloat(alphaInput.value)`. Documented to prevent future confusion.

Tradeoff: BM25 index must be built on Qdrant collection creation. `_ensure_text_index()` is idempotent — safe to call on existing collections. Minimal overhead.

Default: Backend default is pure vector ($\alpha=None$). UI default is $\alpha=0.5$. The separation allows the API to remain backward-compatible while the UI surfaces the recommended setting.

17. Document-Head Extraction Normalizer

Decision: At ingest time, `.htm/.html/.xhtml` files (SEC XBRL, ESEF iXBRL) have a `[Document: <entity> FY<year>]` prefix injected into every chunk before embedding.

Why: Financial filings use anaphora pervasively — "the Company", "the Firm", "the registrant" rather than the entity name. A chunk reading "the Company's CET1 ratio was 15.7%" is semantically indistinguishable from the same sentence in any other filing at embedding time. The prefix resolves anaphora at the embedding layer: "Goldman Sachs FY2024 — the

Company's CET1 ratio was 15.7%" embeds near other Goldman Sachs content, not near generic CET1 content.

Implementation: `pipeline.py` `_extract_document_metadata()` — extracts

`dei:EntityRegistrantName` and `dei:DocumentFiscalYearFocus` from inline XBRL tags.

Fallback chain: authoritative tag → scan → filename stem. Non-markup files (PDF, DOCX, MD) receive null returns — no prefix, no overhead. Controlled by `AISTUDIO_DOC_HEAD_NORMALIZER` env var.

Measured impact: Entity extraction 9/9, year extraction 9/9 on ESEF corpus. SEC 10-K corpus: enables temporal trend queries at K=5 that previously required K=20+ to retrieve both the entity and the relevant year.

Tradeoff: Prefix adds ~8–12 tokens per chunk. Context window budget is slightly reduced. Negligible at current chunk sizes (1,200 chars, ~300 tokens).

18. System Prompt Architecture: File-Based with Hot-Reload

Decision: The system prompt lives in `prompts/system.txt`, loaded at runtime by `api.py` and memoized. A `POST /prompt/reload` endpoint clears the cache, enabling prompt updates without restarting the backend.

Why: Embedding the system prompt in code (`api.py`) ties prompt iteration to deployment cycles. Separating it into a file makes prompt engineering a first-class workflow: edit the file, hit reload, query immediately. No `ais_start` required.

Memoization pattern: `_SYSTEM_PROMPT_BASE` is a module-level variable initialized to `None`. On first query, `api.py` reads and caches the file. `/prompt/reload` sets it back to `None` — next query re-reads from disk.

Prompt design: Citation rules use positive + negative examples, not just ban lists. Small models (llama3.1:8b) respond better to CORRECT/WRONG pairs than to prohibition lists alone — prohibitions suppress the cited behavior but can inadvertently suppress the citation marker entirely. The positive example pattern ("Net revenue increased 12% in fiscal 2024 [2]") anchors the correct output format explicitly.

Tradeoff: System prompt is not version-controlled in git (it's a runtime file, not source code). Tracked via the BUNDLE/PACKET protocol.

19. No Orchestration Framework: LangChain and LlamaIndex Declined

Decision: AIStudio implements retrieval, chunking, prompt assembly and the citation contract directly, against Qdrant and Ollama. No LangChain, no LlamaIndex, no orchestration framework of any kind.

Why: The purpose of this project is substrate knowledge — understanding retrieval well enough to reason about it in an architecture review, not shipping a RAG product fastest. A framework is precisely the wrong tool for that goal. It would have supplied a working pipeline in an afternoon and, in exchange, hidden every decision worth understanding: how chunks are sized and overlapped, how hybrid scores are combined, what actually reaches the model's context window, and why a citation marker appears or does not.

Several findings in this codebase were only reachable because there was no abstraction in the way:

- **Entity anaphora as the dominant retrieval failure** — visible only by reading the chunks a query actually retrieved, which led to the Document-Head normalizer (ADR 17).
- **"Lost in the middle" citation dropout** — diagnosed by controlling context ordering directly and testing a slim system prompt against a verbose one (ADR 18).
- **A memory guard applying a load-time test to an already-resident model** — found by owning the request path end to end. A framework's `.query()` would have returned the same wrong answer with no seam to look through.

None of these are exotic. They are ordinary RAG failure modes, and each was found by reading code we wrote rather than tracing code we imported.

The counter-argument, stated fairly: a framework brings maintained integrations, a large community, and battle-tested edge-case handling. For a team shipping a product against a deadline, that is usually the right trade — and this decision should not be read as a general recommendation against orchestration frameworks. It is a decision about *this* project's purpose.

Tradeoff: Every integration is ours to write and maintain. Adding a vector store, an embedding provider, or a document loader is real work rather than a config change. Accepted deliberately: the integration surface here is small (one vector store, one inference host, two document formats) and the understanding is the deliverable.

Relationship to ADR 1: the same reasoning, one layer down. ADR 1 declines cloud inference and isolates the swap point in `ollama_client.py`; this ADR declines the orchestration layer above it. Both keep the seams visible and under our control. The production path is unchanged — LiteLLM for provider abstraction (ADR 1), and the inference layer remains swappable via the OpenAI-compatible interface described in `architecture_elements.md` §7.

Would change for production: If AIStudio grew a broad integration surface — many vector stores, many providers, many loaders — the maintenance argument would eventually outweigh the transparency one. That threshold is not close.